

CS 480 – INTRODUCTION TO ARTIFICIAL INTELLIGENCE

TOPIC: ADVERSARIAL SEARCH AND GAMES
CHAPTER: 5



Mustafa Bilgic



<http://www.cs.iit.edu/~mbilgic>



<https://twitter.com/bilgicm>

ADVERSARIAL SEARCH

- A **multiagent** and **competitive** environment
- Agents are trying to maximize their utilities at the expense of other agents' utilities
- Zero-sum games, where the sum of the utilities in the end is constant
 - E.g., win = +1, loss = -1, tie = 0
- Games with abstract descriptions attracted the attention of AI researchers quite a bit
 - The states are relatively easy to represent, limited number of actions, precise rules, expected outcomes, etc. E.g., chess, checkers, backgammon
 - Physical games, except soccer, has not attracted much attention

CHAPTER 5 V.S. CHAPTER 3

- Chapter 3:
 - Goal-based agent, where the path to the goal is optimized
- Chapter 5:
 - Utility-based agent, where the expected utility in the end is optimized

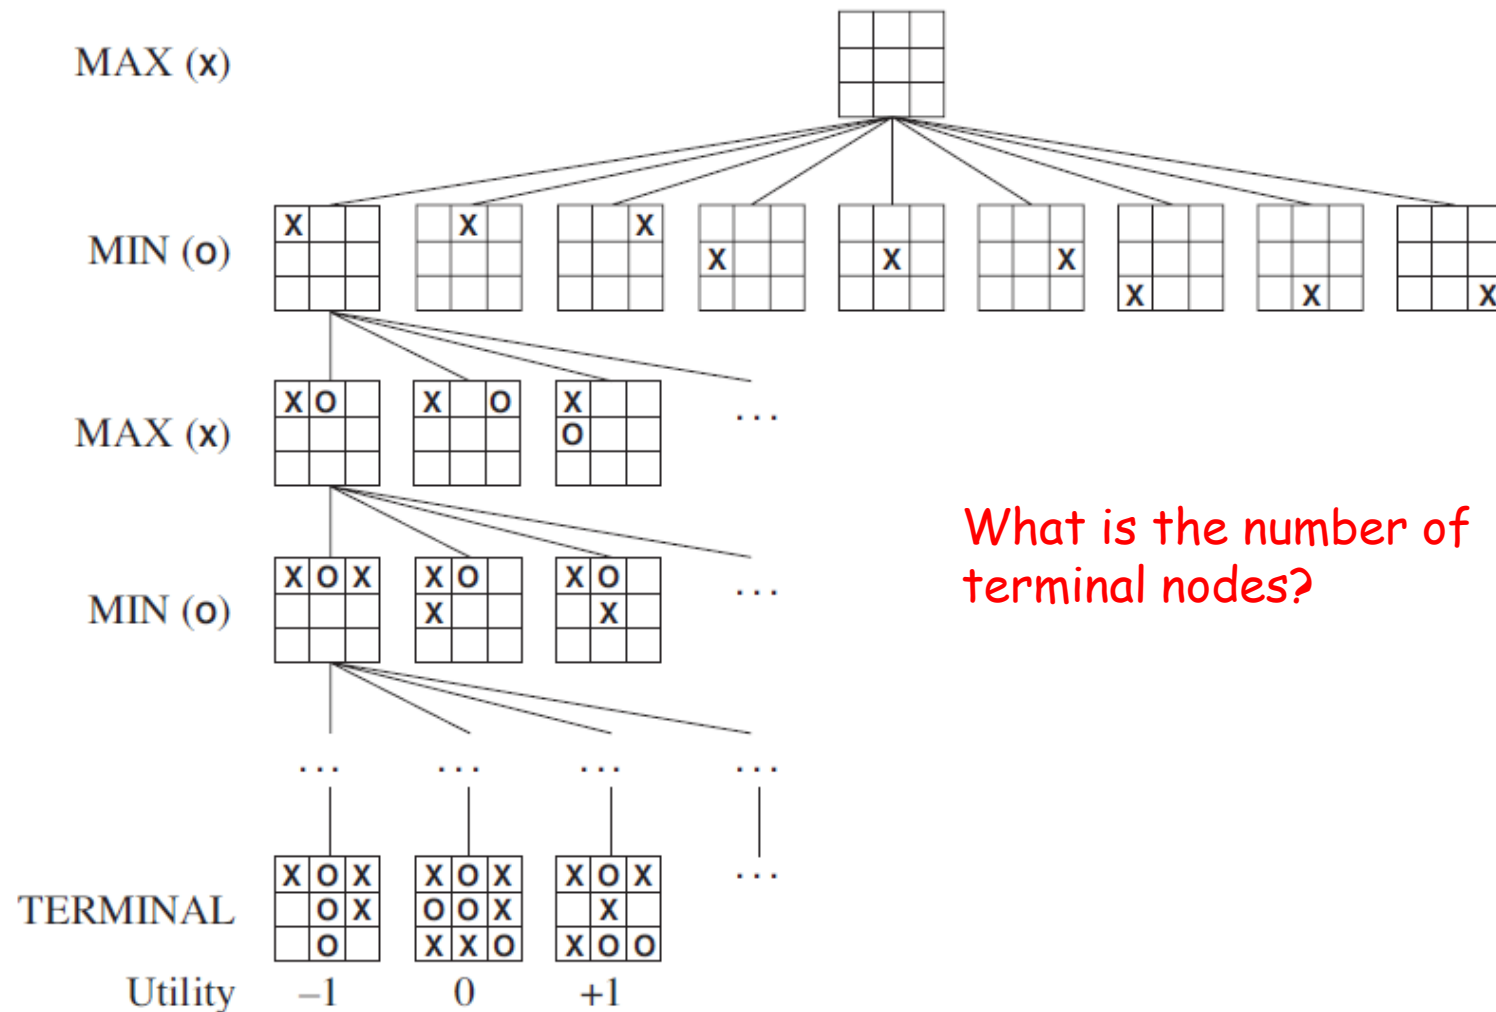
GAMES ARE PRETTY HARD TO SOLVE

- For example, chess
 - Average branching factor: ~ 30
 - If a game lasts 40 moves per player
 - Search tree depth is 80
 - Number of nodes: $30^{80} \approx 10^{118}$
 - Number of estimated atoms in the universe: 10^{80}
 - You still must make a decision before you can calculate the optimal move

GAME DEFINITION

- S_0 : The initial state
- $Player(s)$: Which player's turn in state s
- $Action(s)$: The set of legal moves in state s
- $Result(s, a)$: The transition model
- $Terminal-test(s)$: Is the game over
- $Utility(s, p)$: The utility of player p in state s
- Two players: MAX and MIN. MAX moves first

GAME TREE – TIC-TAC-TOE




OPTIMAL STRATEGY

- Maximize utility
- An optimal strategy leads to outcomes at least as good as any other strategy when one is playing on *infallible* opponent
 - What if the opponent is not playing optimally?
- The **minimax** algorithm

LET'S SEE IT IN ACTION

MAX

MIN




TIME AND SPACE COMPLEXITY

- The minimax algorithm is a depth-first search algorithm
 - Space: $O(bm)$
 - Time: $O(b^m)$
- Can we do better?


ALPHA-BETA PRUNING

- Keep lower and upper bounds (α, β)
- MAX updates the lower bound α
- MIN updates the upper bound β
- Both pass the current bounds to their children
- $\text{MIN}(\alpha, \beta)$ prunes if $v \leq \alpha$; MAX already has a better choice α
- $\text{MAX}(\alpha, \beta)$ prunes if $v \geq \beta$; MIN already has a better choice β

EXAMPLE: MIN PRUNES



EXAMPLE: MAX PRUNES



ALPHA-BETA PRUNING ALGORITHM

function ALPHA-BETA-SEARCH($state$) **returns** an action
 $v \leftarrow \text{MAX-VALUE}(state, -\infty, +\infty)$
 return the *action* in $\text{ACTIONS}(state)$ with value v


function MAX-VALUE($state, \alpha, \beta$) **returns** a *utility value*
 if $\text{TERMINAL-TEST}(state)$ **then return** $\text{UTILITY}(state)$
 $v \leftarrow -\infty$
 for each a **in** $\text{ACTIONS}(state)$ **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \geq \beta$ **then return** v
 $\alpha \leftarrow \text{MAX}(\alpha, v)$
 return v

function MIN-VALUE($state, \alpha, \beta$) **returns** a *utility value*
 if $\text{TERMINAL-TEST}(state)$ **then return** $\text{UTILITY}(state)$
 $v \leftarrow +\infty$
 for each a **in** $\text{ACTIONS}(state)$ **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$
 if $v \leq \alpha$ **then return** v
 $\beta \leftarrow \text{MIN}(\beta, v)$
 return v


ALPHA-BETA PRUNING

MAX

MIN




ALPHA-BETA PRUNING – ANOTHER EXAMPLE



ALPHA-BETA PRUNING (FIGURE 5.7)

MAX

MIN




Can we do better?

MOVE ORDERING

MAX

MIN



What if we re-order these?

STILL

- To be able to prune, the alpha-beta algorithm has to visit some of the leaf nodes
 - Of course, this is not practical for most games
- A solution: early cut-off
 - How can we do this? Remember, the utility function is defined for the terminal states.


HEURISTIC EVALUATION FUNCTIONS

- An estimate of the utility is returned rather than the actual utility
- Define features and a weighted feature combination function
 - e.g., $\text{EVAL}(s) = w_1f_1(s) + w_2f_2(s) + \dots + w_nf_n(s)$
- Chess
 - Pawn: 1, Knight or bishop: 3, rook: 5, queen: 9

HEURISTIC EVALUATION FUNCTIONS

- Just like in Chapter 3, heuristic functions are not part of the problem description
 - It takes additional expertise and learning to devise useful heuristic functions
- They are not very easy to define

LINEAR COMBINATION FAILS



(a) White to move



(b) White to move

CUTTING-OFF IS NOT TRIVIAL

- A simple depth limit will not work
 - Cut-off at positions that are quiescent
 - Positions in which a drastic change is unlikely to happen
- Horizon effect
 - A serious damage is unavoidable but the search delays it by making little sacrifices
 - The little sacrifices are unnecessary as the serious damage is unavoidable but a depth-limited search cannot see it

HORIZON EFFECT




Figure 5.9 The horizon effect. With Black to move, the black bishop is surely doomed. But Black can forestall that event by checking the white king with its pawns, encouraging the king to capture the pawns. This pushes the inevitable loss of the bishop over the horizon, and thus the pawn sacrifices are seen by the search algorithm as good moves rather than bad ones.

SEARCH V.S. LOOKUP

- Opening of a game
 - A depth-limited search is unlikely to help
 - Human expertise and learning from previous games for opening moves
- Ending a game
 - When very few pieces are left, the number of possible states is not many; rather than a depth-limited search, look-up a winning strategy

Monte-Carlo Tree Search

```
function MONTE-CARLO-TREE-SEARCH(state) returns an action
  tree  $\leftarrow$  NODE(state)
  while IS-TIME-REMAINING() do
    leaf  $\leftarrow$  SELECT(tree)
    child  $\leftarrow$  EXPAND(leaf)
    result  $\leftarrow$  SIMULATE(child)
    BACK-PROPAGATE(result, child)
  return the move in ACTIONS(state) whose node has highest number of playouts
```

Figure 5.11 The Monte Carlo tree search algorithm. A game tree, *tree*, is initialized, and then we repeat a cycle of SELECT / EXPAND / SIMULATE / BACK-PROPAGATE until we run out of time, and return the move that led to the node with the highest number of playouts.

Monte-Carlo Tree Search




Figure 5.10 One iteration of the process of choosing a move with Monte Carlo tree search (MCTS) using the upper confidence bounds applied to trees (UCT) selection metric, shown after 100 iterations have already been done. In (a) we select moves, all the way down the tree, ending at the leaf node marked 27/35 (for 27 wins for black out of 35 playouts). In (b) we expand the selected node and do a simulation (playout), which ends in a win for black. In (c), the results of the simulation are back-propagated up the tree.

SO FAR

- Deterministic
- Fully-observable

STOCHASTIC GAMES

- An element of chance, such as dice
- Modify the search tree and the minimax algorithm:
 - Include a chance node, that leads to all possible outcomes
 - Rather than computing utilities, compute expected utilities

PARTIALLY OBSERVABLE GAMES

- The current state is only partially observable
- Model what is not observed
- Examples
 - Card games
 - Battleships
 - Kriegspiel

HISTORICAL NOTES